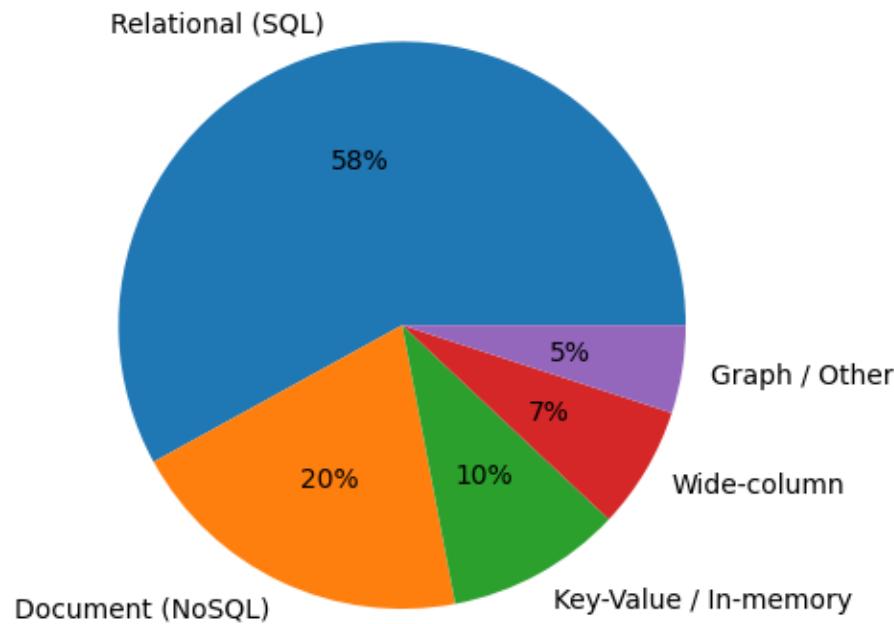


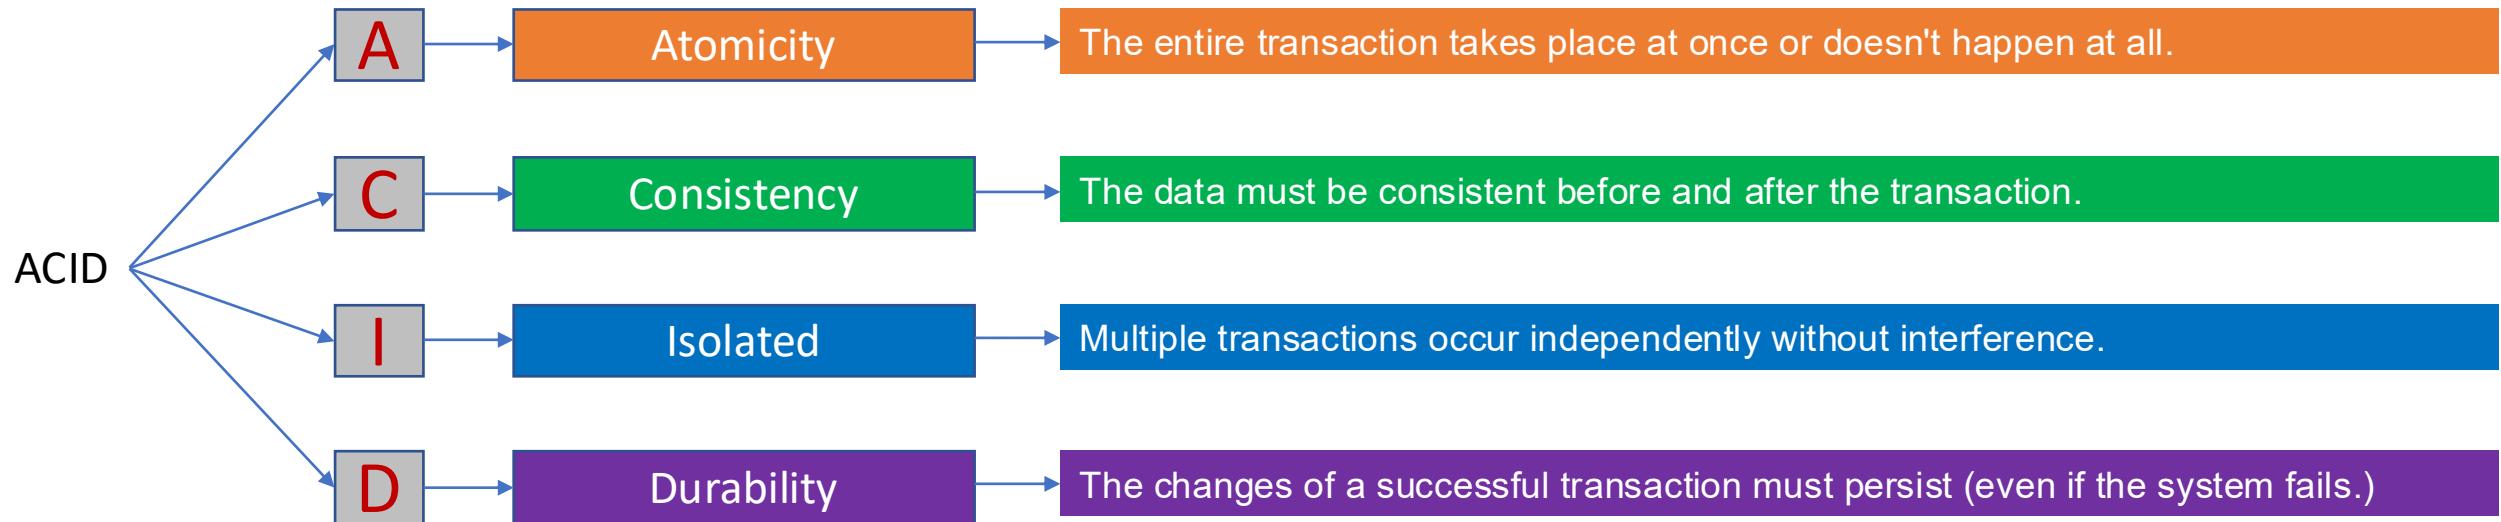
# Big data management: CAP theorem and NoSQL

# Relational database technology

- Invented in the 1970s
- Well-founded theoretically (relational algebra)
- Provides **high-consistency**
- Many added capabilities
  - BLOB storage
  - Document store
  - Text search
  - In-database code execution
  - Data warehousing (star schema)
  - GIS extensions
- Excellent to cope with **centralised storage**

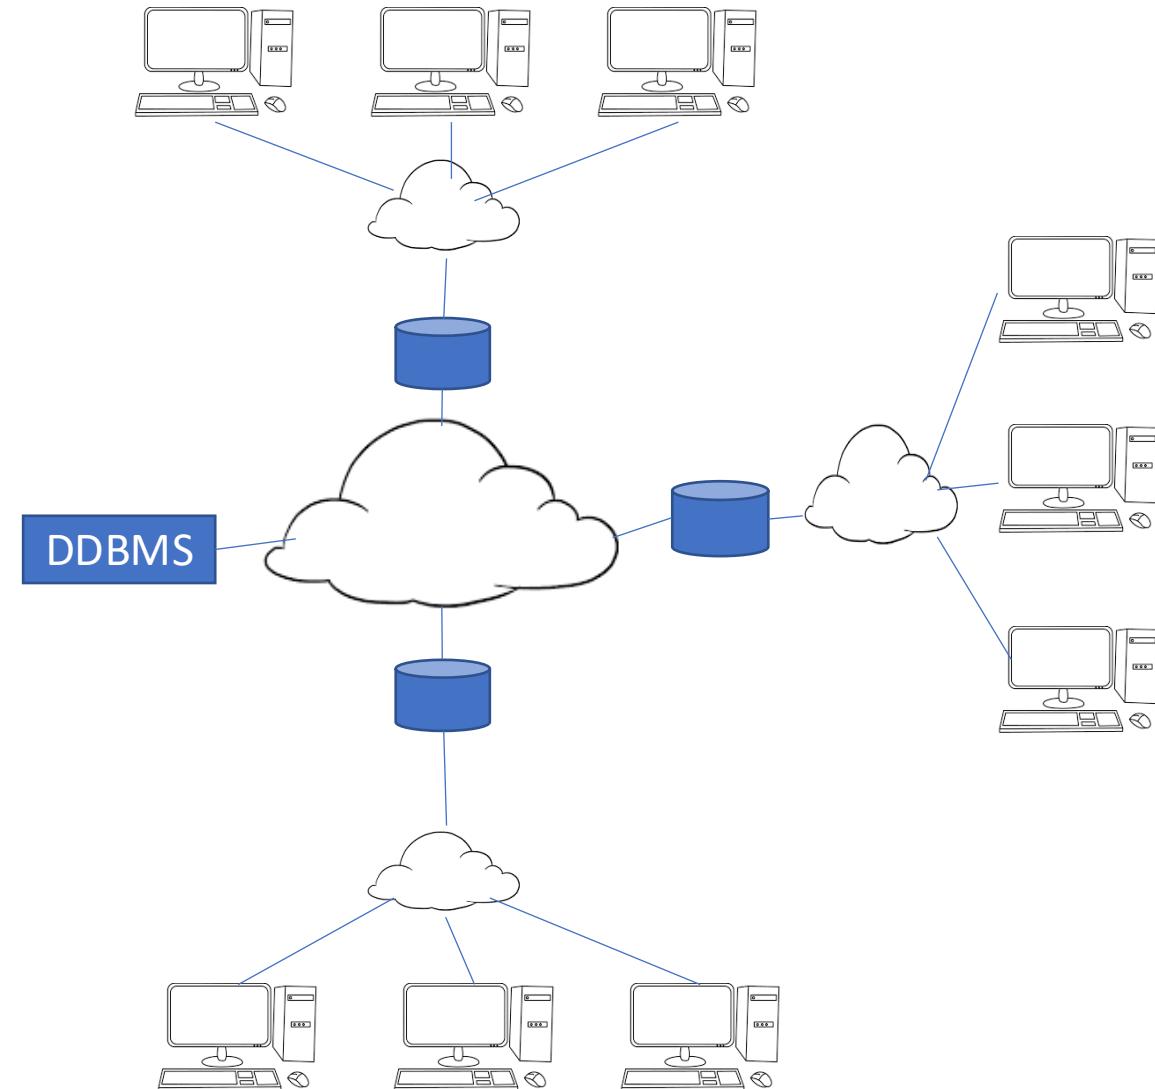


# ACID properties in RDBMS



# RDB technology and Internet

- Intensive online access to data
- Requires
  - Availability, scalability, and speed
- Consistency cannot be ensured in all times

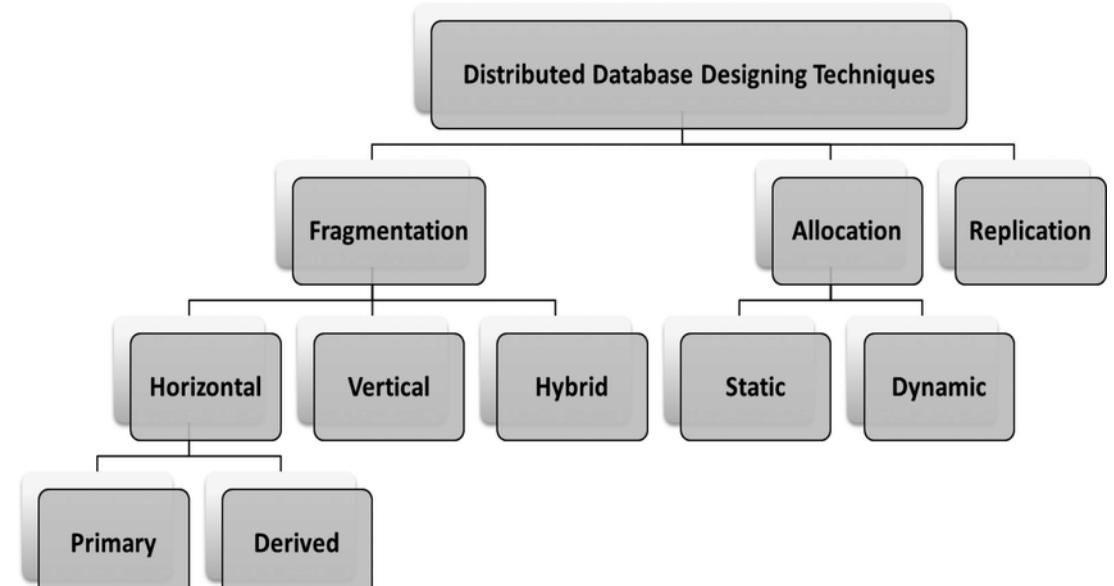
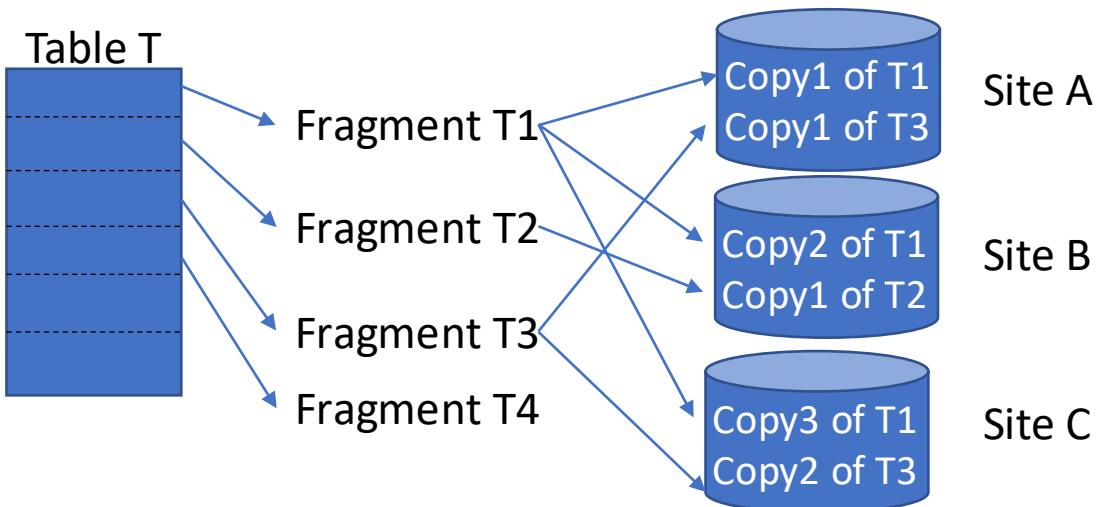


# Limitations of relational databases

- Low response time with massive amounts of data
  - Fast querying difficult
- Expensive Scale up
  - Adding memory, CPU...
- Limited analytical capabilities
  - Need OLAP like schemas
- Appearance of distributed databases
  - Data stored in different nodes
  - Nodes are interconnected by a computer network
  - Nodes are databases
  - DDBMS software manages DDB (access, transparent distribution)

# DDBMS: Fragmentation (or partitioning) and Replication

- User unaware of **fragments** and their **replication**
- Queries specified on the relations (original schema) and not on fragments



# Example partitioning

## Horizontal

ID	Name
1	Shaun
2	Tao
3	Ray
4	Jesse
5	Robin

ID	Name
4	Jesse
5	Robin

shaun @ geekswithblogs

## Vertical

ID	Name	Avatar
1	Shaun	<Binaries>
2	Tao	<Binaries>
3	Ray	<Binaries>
4	Jesse	<Binaries>
5	Robin	<Binaries>

ID	Name
1	Shaun
2	Tao
3	Ray
4	Jesse
5	Robin

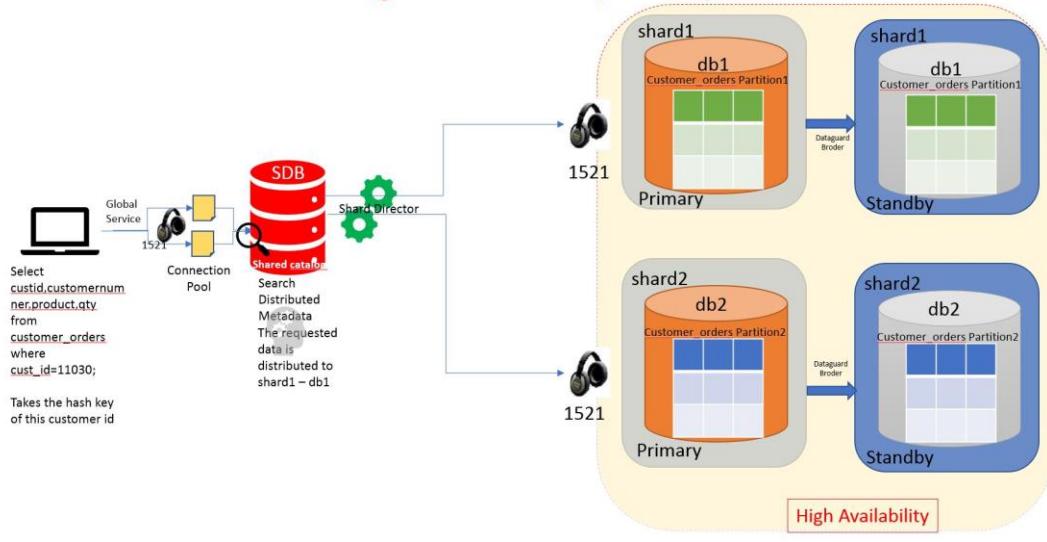
ID	Avatar
1	<Binaries>
2	<Binaries>
3	<Binaries>
4	<Binaries>
5	<Binaries>

shaun @ geekswithblogs

# Database sharding

- Another way of naming **partitioning**
- Table portioned into **shards**
- Shards are necessarily stored on different computers
- Queries are run on separate shards

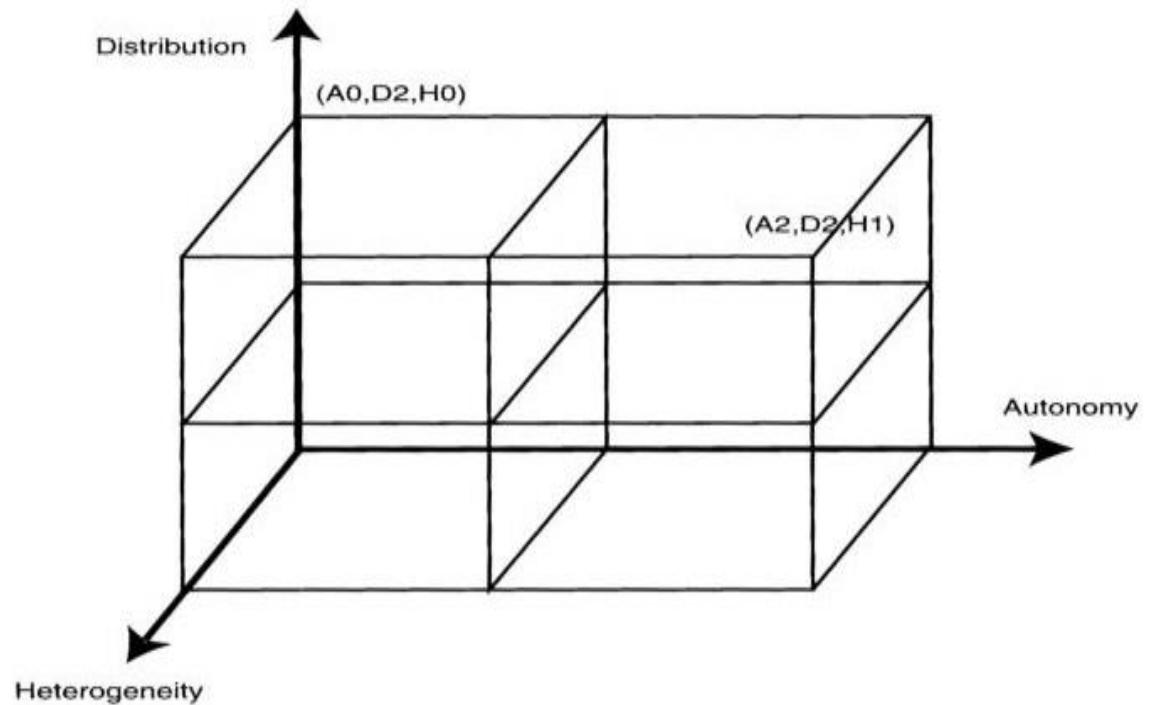
Oracle Sharding Architecture (12c R2)



In 2011, Facebook split its MySQL database into **4,000 shards** in order to handle the site's massive data volume

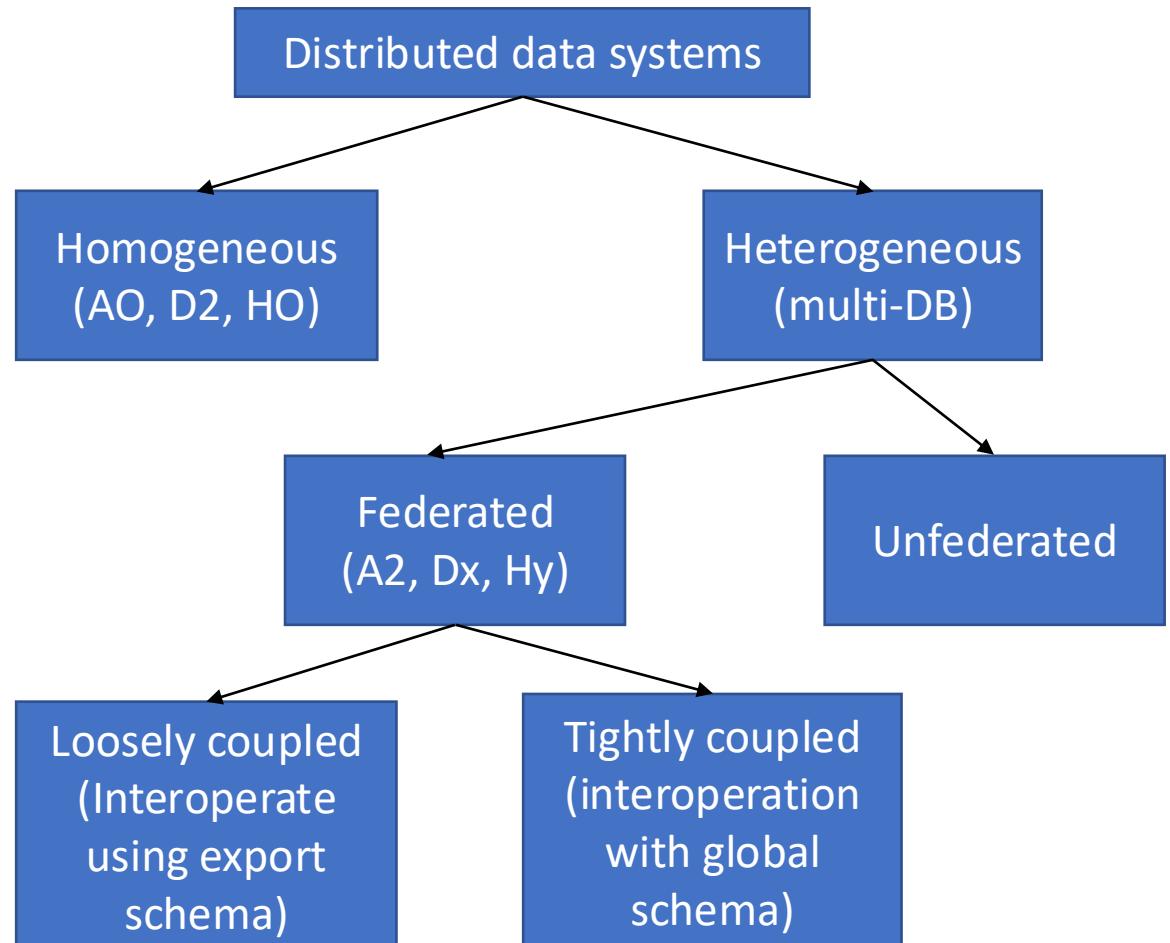
# Dimensions of distributed database systems

- Distribution
  - What is the mode of data sharing?
  - D<sub>0</sub>: No-distribution
  - D<sub>1</sub>: client-server
  - D<sub>2</sub>: peer-to-peer
- Heterogeneity
  - Do local systems use same or different technologies?
  - H<sub>0</sub>: Homogeneous
  - H<sub>1</sub>: Heterogeneous
- Autonomy
  - How much local nodes operate independently ?
  - A<sub>0</sub>: tight integration, a single coordinator
  - A<sub>1</sub>: semi-autonomous, but implement sharing
  - A<sub>2</sub>: independent, sharing done by another software



# Taxonomy of distributed database systems

- Distributed database is a logically integrated collection of shared data physical distributed across network nodes
- In DBMS, global schema is union of local schemas of all databases
- In multi-DB, global schema is a subset of the union of all local databases



# Challenges of Distributed Database systems

- Large-scale internet applications require scalable storage and support of concurrency
  - Increasesable data capacity
  - Growing read/write throughput
  - Handle highly concurrent access
- Horizontal scale required to allow for throughput and capacity increases, however
  - Database system must maintain data consistency while read/write operations are conducted
  - In case of node failure, the overall system must maintain availability
  - Operations need low latency

# Properties of distributed systems

- **Consistency**

- The system assures that **operations are atomic**, and changes are disseminated simultaneously to all nodes, with same results
- A read is guaranteed to return the most recent write

- **Availability**

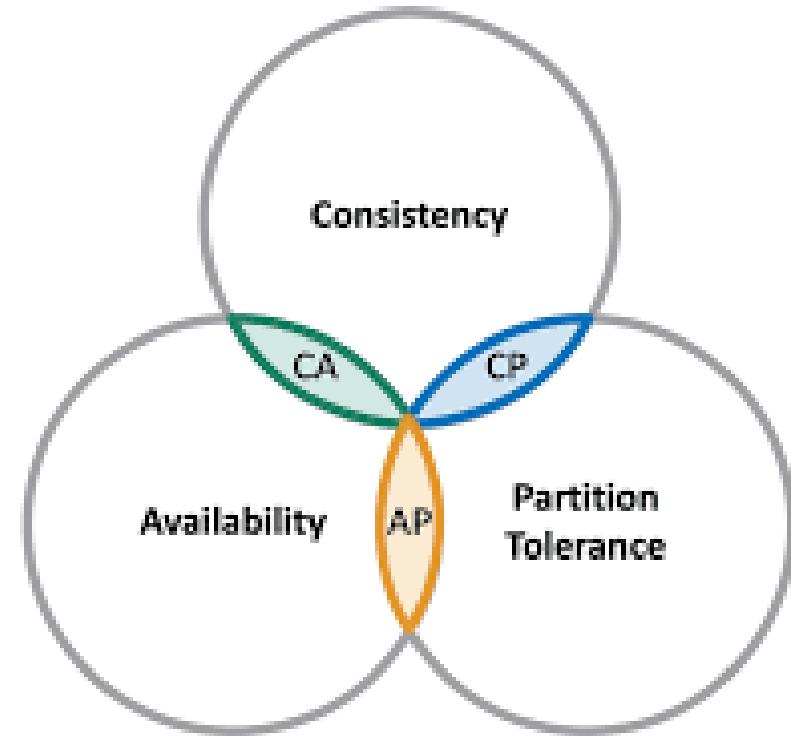
- The system must answer **every request**, even in case of failures
- A non-failing node returns a reasonable response within a reasonable time, **without guarantee of being the latest write**

- **Partition tolerance**

- Partition is any arbitrary **split between nodes** resulting in message loss in between
- The system must be **resilient to message losses between nodes**
  - Consistency mechanisms must cope with message losses
  - Availability requires any partition to be able to answer a request

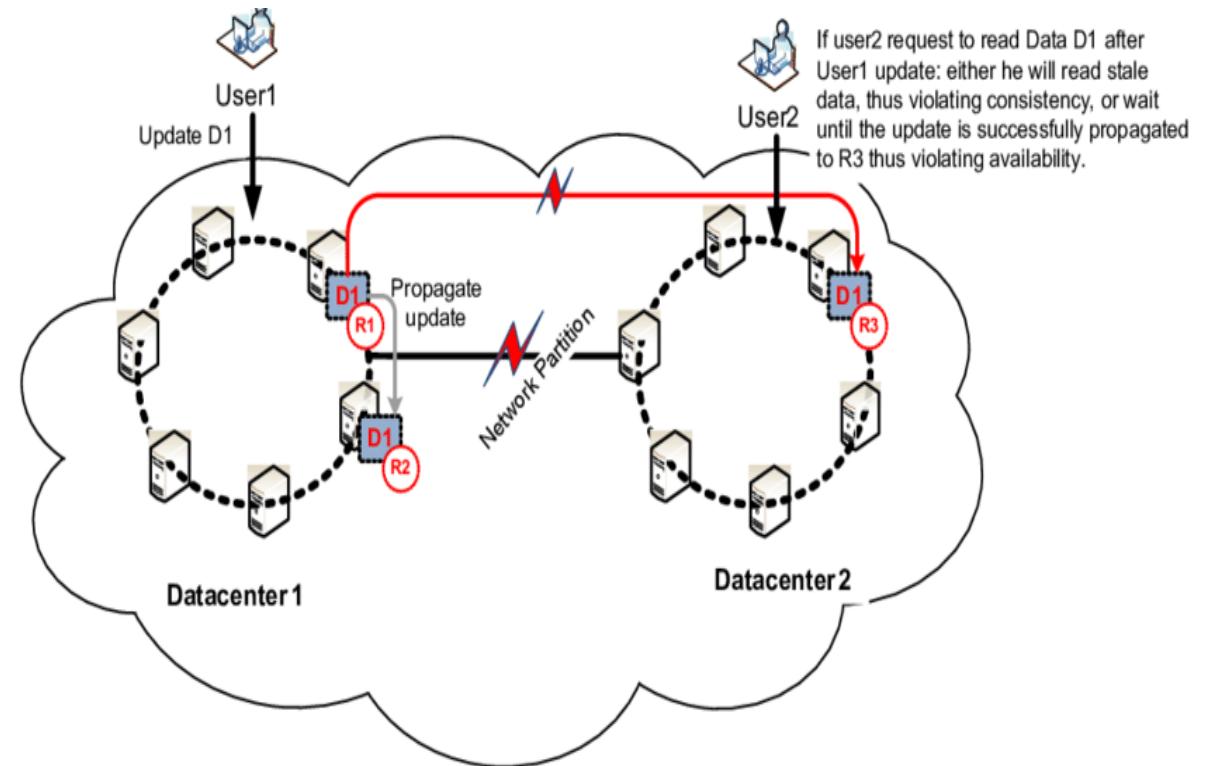
# CAP theorem

- *A distributed system can only have at most two of the three properties:*
  - CA, CP, AP
- This was initially a **Conjecture** by Brewer 2000
- It has been demonstrated as a **Theorem** by Gilbert and Lynch 2002



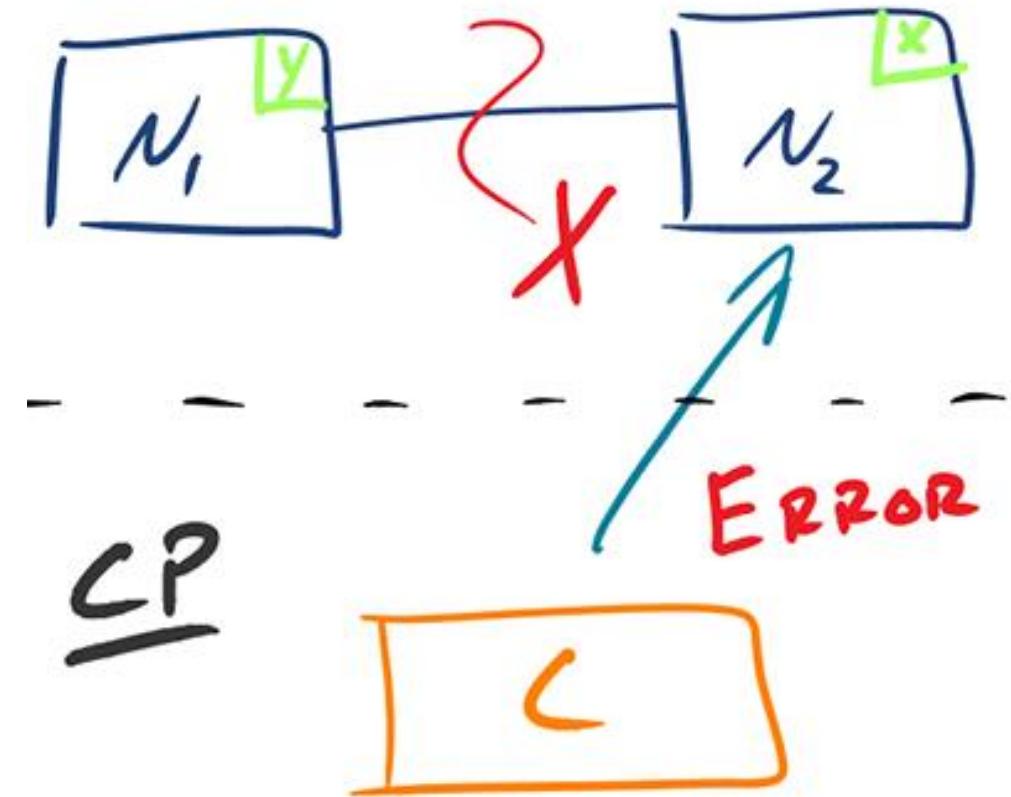
# Consistency and Availability (CA)

- Such systems provide consistent and available services
- **Do not tolerate partitions**
  - In case of partition, systems become inconsistent
- **RDMS follow this approach**
  - Replication is used to achieve high availability
  - Transaction protocols (e.g. two-phase commit) ensure consistency
  - Partitions lead to conflicting replicas
- Two options to recover from such situations
  - **Favour consistency (CP)**
    - require consensus protocol
    - No answer to request if no consensus
    - sacrifice availability
  - **Favour availability (AP)**
    - Returned data is not guaranteed to be the latest update



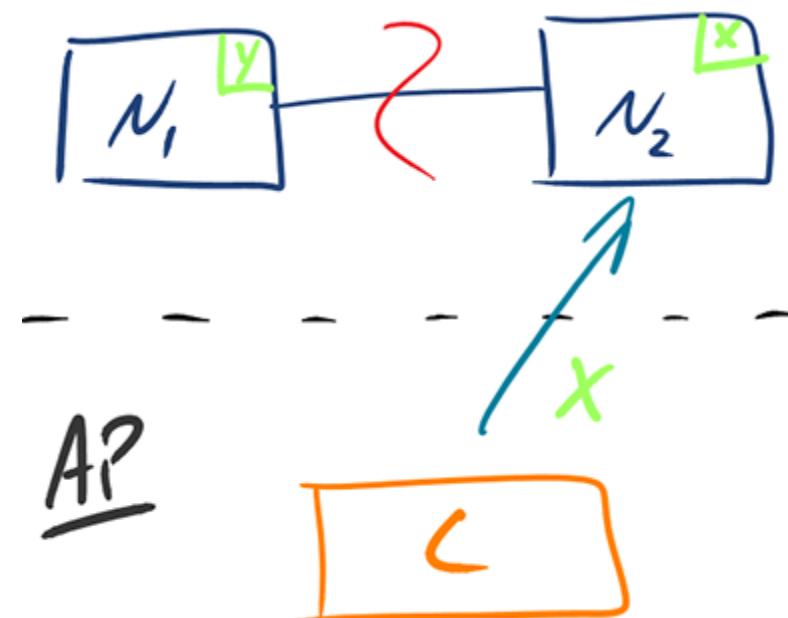
# Consistency and Partition tolerance (CP)

- Systems with CP model provide **strongly consistent** service, even in the presence of partition
- Nodes must agree to ensure consistency
  - Use quorum and majority decision algorithms (e.g. Paxos protocol)
- In case of **partition**, consensus is delayed, therefore **affecting availability** (no answer possible)



# Availability and Partition tolerance (AP)

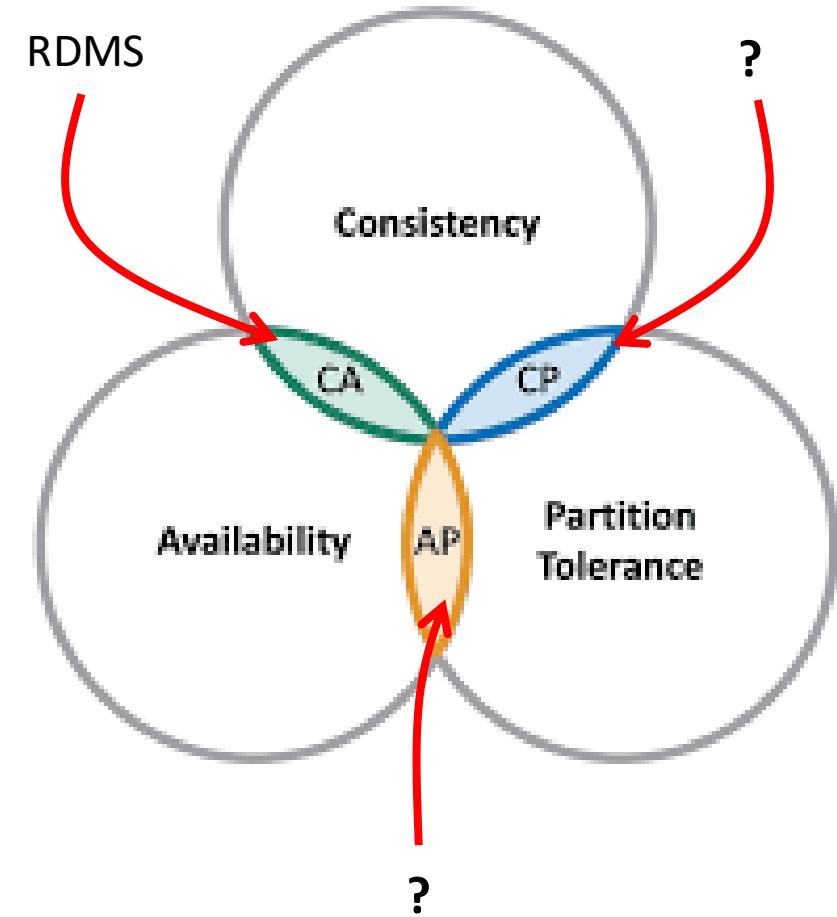
- Systems with AP model provide **available service**, even in the presence of partition
- Unreachable nodes hold **temporarily inconsistent** data
- **Eventual consistency** consists in ensuring the system converges to full consistency



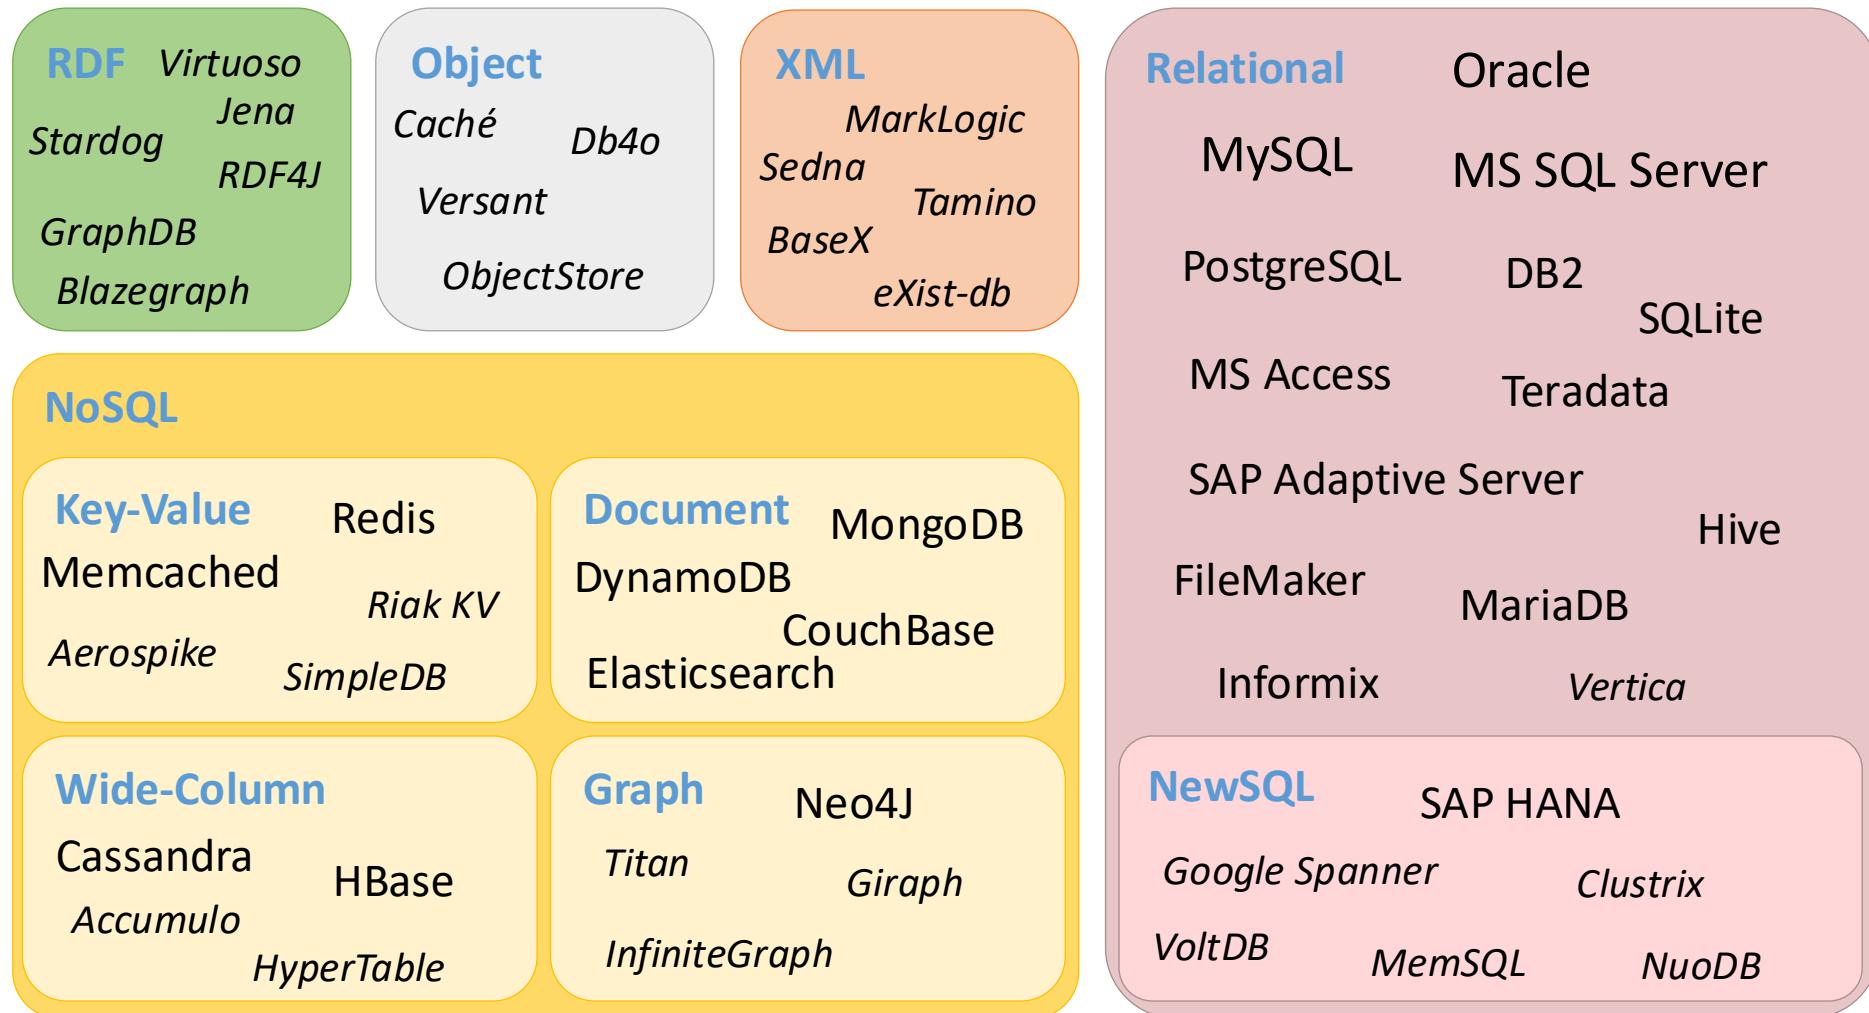
# Question

- Since RDMS have CA model, therefore cannot provide availability or consistency in the presence of partition
  - Which database systems can be designed according to CP or AP models?

NoSQL database systems



# Many database systems



# noSQL movement: history

- noSQL used first by Carlo Strozzi to query a RDB with no SQL (1998)
- NoSQL arbitrary name of a conference hosted by Rackspace (2009)
  - Presenting a collection of distributed non-relational databases
- Started as “no more SQL”, an anti SQL movement
- Became “not only SQL”, gentler definition
- Driven by the web
  - Minimally structured databases
  - Scale large amount of data across servers

2000



2004



2005



2007

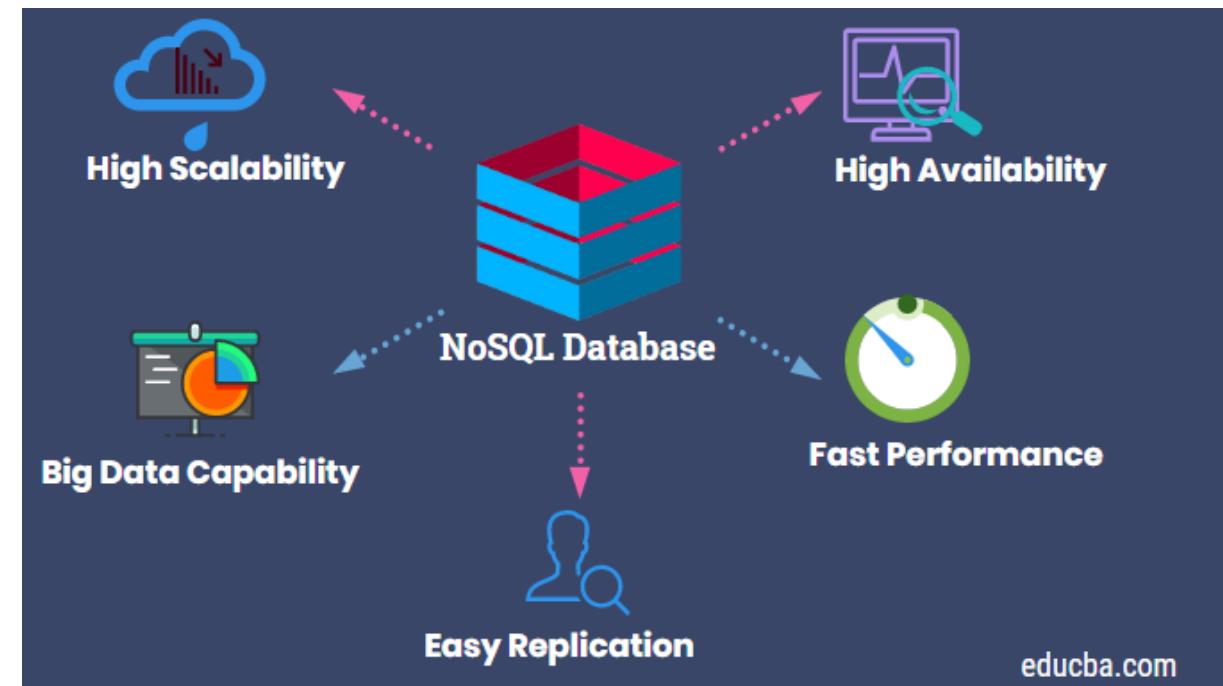


2008



# Characteristics of noSQL databases

- Do not respect the ACID principles
  - Consistency cannot be achieved while targeting high availability
  - Substituted by **eventual consistency** (in a parallel environment)
- No standard declarative query language
  - Most use variations of SQL
    - Cosmos DB, Cassandra CQL, Elasticsearch
    - SQL, Cockroach Labs, MongoDB
  - Couchbase uses N1QL (SQL like)
- Do not have predefined schema
  - Rely on dynamic schema
- Built to scale horizontally (vs. vertical scalability of RDBMS)



# BASE principles

- Basic availability

- Highly distributed DBM, with replication, to provide availability even with failures

- Soft State

- Consistency is the problem of the developer not the DBMS

- Eventual consistency

- Database should converge to consistent state in the future (operations can be executed without waiting for prior consistency)

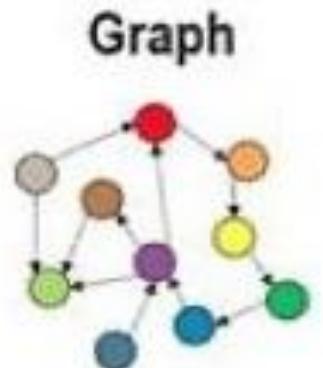
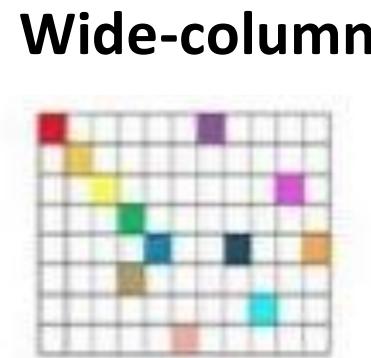
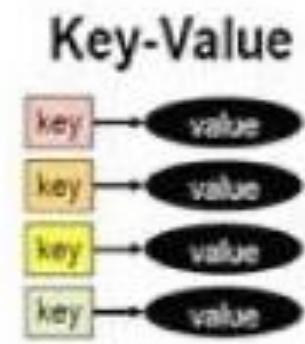
## ACID

- Strong Consistency
- Isolation
- Focus on "commit"
- Nested transactions
- Less Availability
- Conservative (pessimistic)
- Difficult evolution (e.g. schema)

## BASE

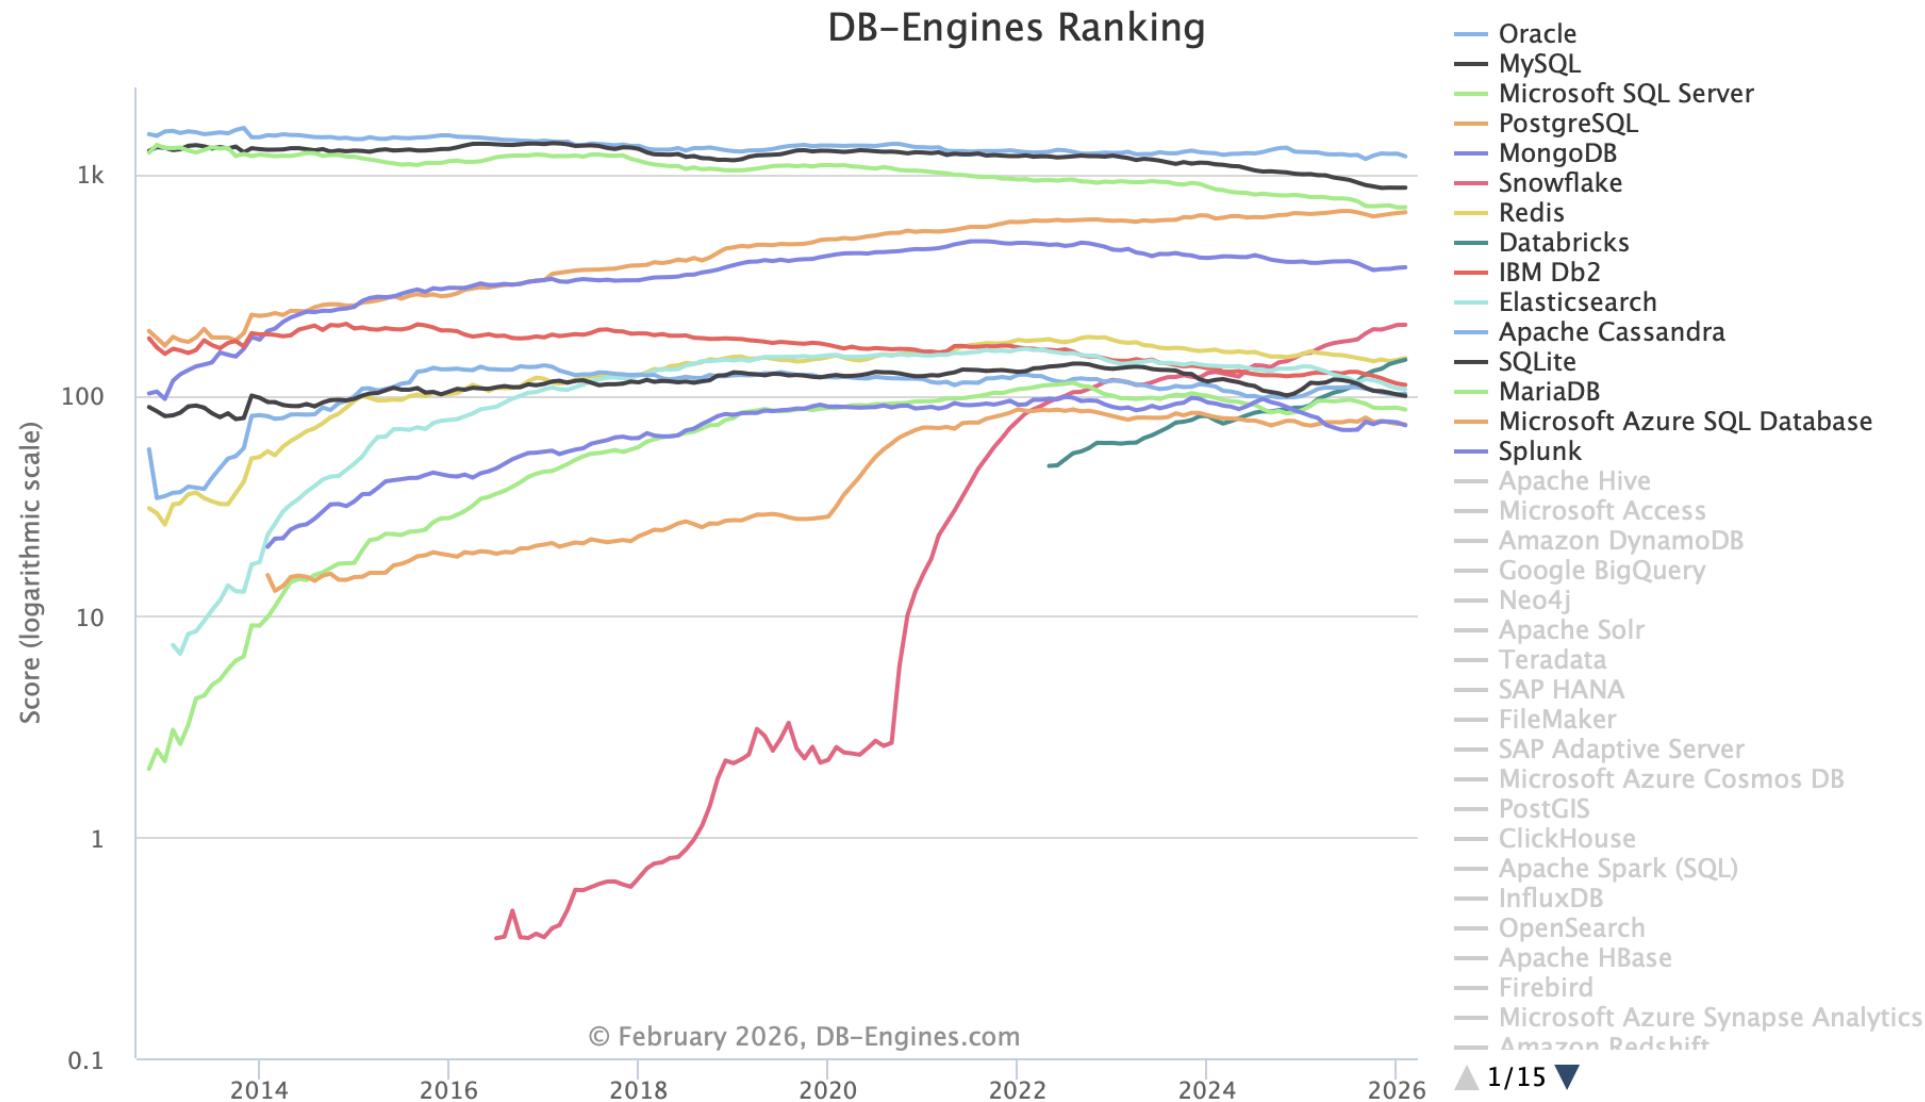
- Weak Consistency
- Availability first
- Best effort
- Approximated answers
- Aggressive(optimistic)
- Simpler!
- Faster
- Easier evolution

# Four noSQL data models



<b>Key-Value Store</b>	 redis	 riak
<b>Wide Column Store</b>	 HBASE	 cassandra
<b>Document Store</b>	 mongoDB	 CouchDB relax
<b>Graph Store</b>	 Neo4j	 InfiniteGraph The Distributed Graph Database

# Database systems usage



# Key-value

- Based on the **key-value pattern**
- Data is stored in **hashtable** (dictionary or map) data structure as objects
- Objects are **opaque** to the database (uniquely read written)
- Objects are **indexed by a key**
- Schema-less
- Mostly in-memory



**ORACLE®**  
NOSQL DATABASE



**amazon**  
DynamoDB  
(also document)



Azure Cosmos DB  
(also document)

# Document

- Extend key-value model to key-document lookup
- Data is structured (set of key-value) and typically stored in JSON or XML
  - Hence the term document
- Documents are indexed by unique key
- Database has access to internal structure of document therefore provide access to fields...
- Indexes are usually used for fast query
- Apache CouchDB and MongoDB use JSON markup in a document data store
- Include search capabilities (search patterns)



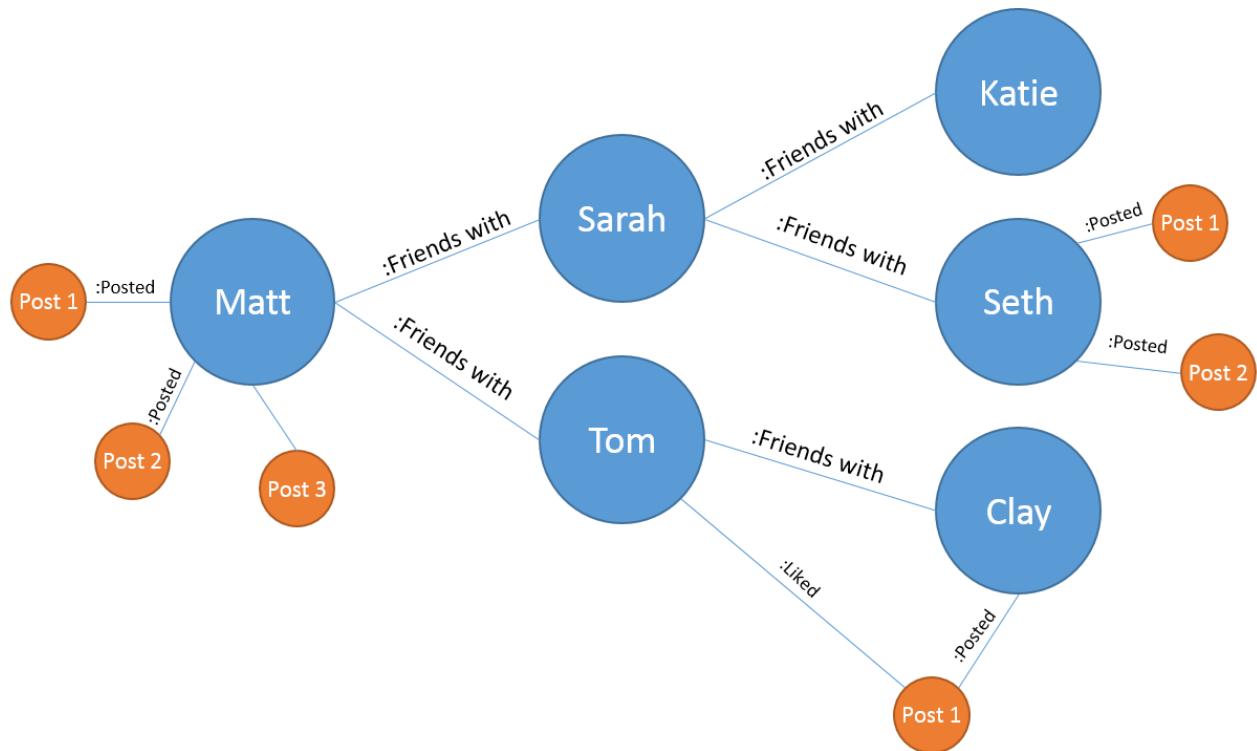
# Wide-column

- Also called **column family databases**
- Use “**flexible**” tables (rows and columns)
  - Sparse data matrix
- Names and formats of **columns** can **vary** from a row to another
  - Not all columns have values
- **Vertical sharding**
  - Column families (set of columns) stored on separate computers
- **Horizontal sharding**
  - Within a column-family , rows are stored on different computers

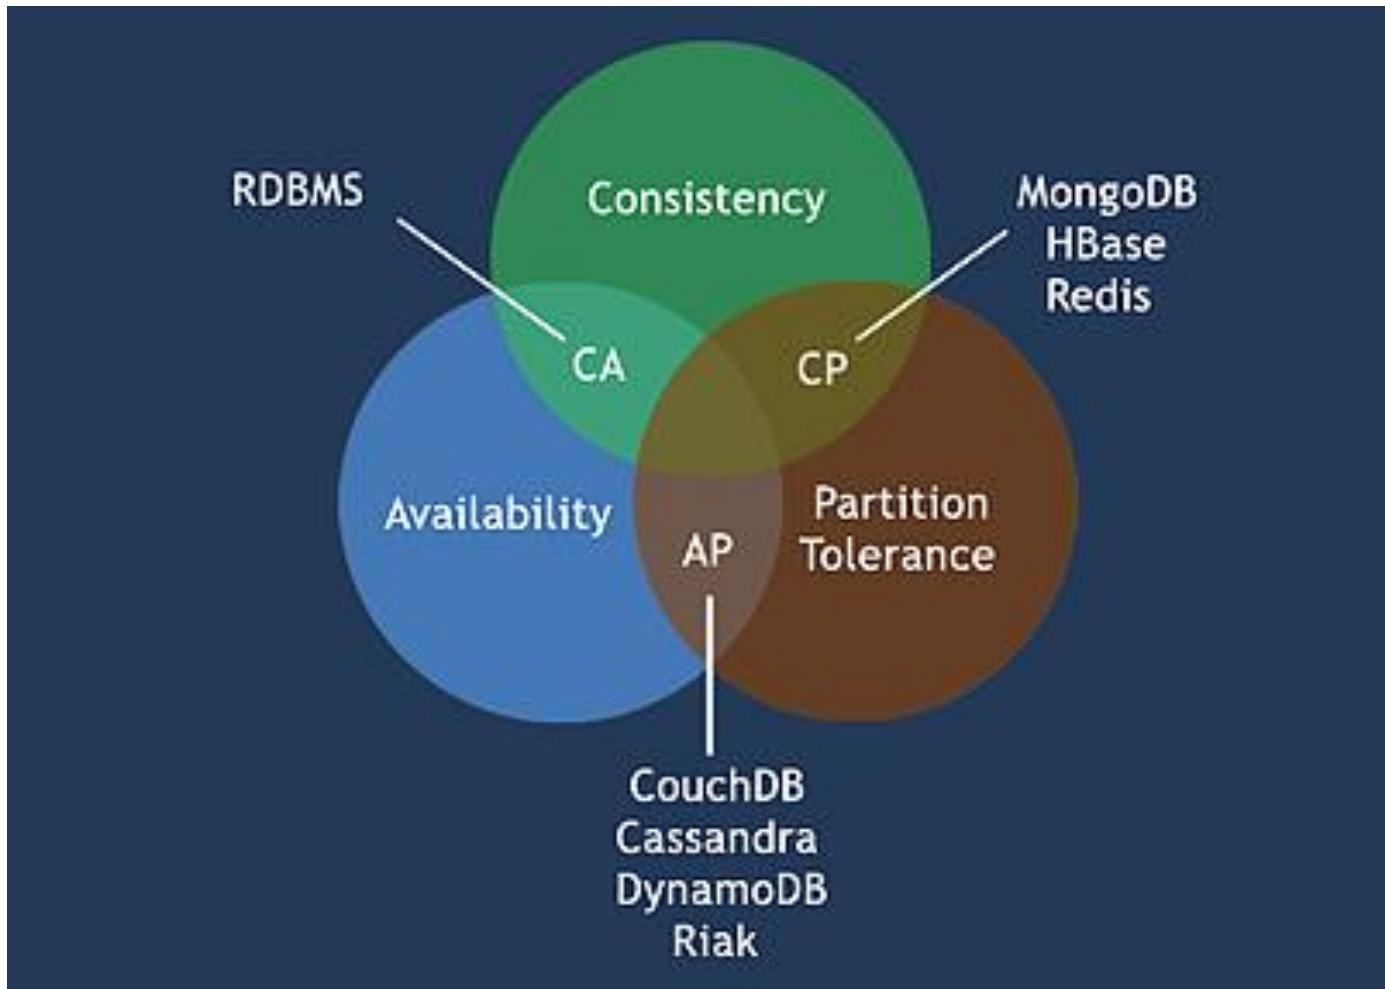


# Graph

- Focus on **relationship between entities**
  - **Nodes** are entities
  - **Edges** are relations (properties)
- Widely used by social networks



# NoSQL and CAP properties



# Critics of CAP theorem

- Trade-off between Consistency and Availability in case of partition is obvious
- Even in the absence of partition (normal operations), one has to choose between **Latency** and **Consistency**
- High availability requires **replication**
- Data updates sent to all replicas simultaneously
  - Without pre-processing layer -> Latency
  - With pre-processing layer -> consistency
- Data updates sent to agreed-upon location first
  - Synchronous -> consistency
  - Asynchronous -> latency (e.g. PNUTS)
- Data updates sent to an arbitrary location (e.g. Dynamo, Cassandra, Riak)
  - Synchronous -> consistency
  - Asynchronous -> latency

# PACELEC theorem

- on **Partition** tradeoff between **Availability** and **Consistency**,
- **Else** tradeoff between **Latency** and **Consistency**

DDBMS	PA	PC	EL	EC	?
BigTable/HBase		✓		✓	
Cassandra	✓			✓	
Cosmos DB	✓			✓	
Couchbase		✓	✓		✓
DynamoDB	✓		✓		
FaunaDB		✓	✓		✓
Hazelcast IMDG	✓	✓	✓		✓
Megastore		✓			✓
MongoDB		✓		✓	
MySQL Cluster		✓			✓
PNUTS		✓	✓		
Riak	✓		✓		
VoltDB/H-Store		✓		✓	

# NoSQL aims

- Flexible
  - no fixed schema
- Highly available
  - replication across compute clusters
- Web-scale
  - automatic horizontal sharding

# NoSQL do not support

- Schema: Data model can evolve
  - Data typically exchanged as JSON
- SQL: No standard query language
  - Proliferation of APIs
  - Simpler lighter-weight interactions
- Transaction processing: no ACID
  - BASE consistency: Basically Available, Soft state, Eventually consistent

# NoSQL characteristics

- **Distributed!**
  - Sharding: splitting data over servers “horizontally”
  - Replication
- **Lower-level** than RDBMS/SQL
  - Simpler ad hoc APIs
  - Programmer ensures consistency (programming not querying)
  - Operations simple and cheap
- **Different flavours** (for different scenarios)
  - Different CAP emphasis
  - Different scalability profiles
  - Different query functionality
  - Different data models



Not  
Only SQL

# RDB vs. noSQL

- Relational Databases don't solve everything
  - SQL and ACID add overhead
  - Distribution not so easy
- NoSQL: what if you don't need SQL or ACID?
  - Something simpler
  - Something more scalable
  - Trade efficiency against guarantees

# Query language standardization

